

Source: advanced_features.md

Advanced Iceberg Features

IceFrame provides a suite of advanced features to give you a complete Iceberg experience, bridging gaps in the underlying libraries.

Iceberg Views

Manage cross-engine views (if supported by your catalog).

```
# Create a view
sql = "SELECT * FROM source_table WHERE id > 100"
ice.create_view("my_view", sql, replace=True)

# Drop a view
ice.drop_view("my_view")
```

Advanced Compaction

Optimize your data layout for better query performance.

```
from iceframe.compaction import CompactionManager

table = ice.get_table("my_table")
compactor = CompactionManager(table)

# Bin-pack small files into 128MB files
stats = compactor.bin_pack(target_file_size_mb=128)

# Sort data (Z-order approximation)
stats = compactor.sort(sort_order=["region", "ts"])
```

Partition Evolution

Evolve your table partitioning without rewriting data.

```
# Get evolution helper
evolver = ice.evolve_partition("my_table")

# Add partitions
evolver.add_day_partition("created_at")
evolver.add_bucket_partition("user_id", num_buckets=16)

# Remove partition
evolver.remove_partition("region")
```

Stored Procedures

Execute maintenance tasks using a familiar procedure call interface.

```
# Rewrite data files (compaction)
ice.call_procedure("my_table", "rewrite_data_files", target_file_size_mb=256)

# Expire snapshots (cleanup)
ice.call_procedure("my_table", "expire_snapshots", older_than_ms=...)

# Remove orphan files (GC)
ice.call_procedure("my_table", "remove_orphan_files")

# Fast-forward branch (WAP)
ice.call_procedure("my_table", "fast_forward", to_branch="audit_branch")
```

Merge-on-Read (MoR) Support

Note: Currently limited by underlying library support.

IceFrame includes structure for MoR writers (`MoRWriter`) to handle Position Deletes and Equality Deletes as support becomes available.

Garbage Collection

Parallelized cleanup operations for large tables.

```
from iceframe.gc import GarbageCollector

gc = GarbageCollector(table)
gc.expire_snapshots(retain_last=5)
```

Source: ai_agent.md

AI Agent

IceFrame includes an AI agent that provides a natural language interface for interacting with your Iceberg tables.

Features

Natural Language Queries: Ask questions in plain English

Schema Discovery: Explore tables and understand data structures

Code Generation: Get Python code for complex operations

Query Optimization: Receive suggestions for better performance

Multiple LLM Support: Works with OpenAI, Anthropic Claude, or Google Gemini

Setup

1. Install Dependencies

```
pip install "iceframe[agent]"
```

This installs: `openai`, `anthropic`, `google-generativeai`, and `rich`.

2. Configure LLM

Set one of these environment variables:

```
# OpenAI (GPT-4, GPT-3.5-turbo)
export OPENAI_API_KEY="your-key"

# Anthropic (Claude)
export ANTHROPIC_API_KEY="your-key"

# Google Gemini
export GOOGLE_API_KEY="your-key"
```

Optional: Specify provider and model explicitly:

```
export ICEFRAME_LLM_PROVIDER="openai" # or "anthropic" or "gemini"
export ICEFRAME_LLM_MODEL="gpt-4"      # or specific model
```

Python API

```
from iceframe import IceFrame
from iceframe.agent.core import IceFrameAgent

# Initialize
ice = IceFrame(config)
agent = IceFrameAgent(ice)

# Chat
response = agent.chat("What tables do I have?")
print(response)

response = agent.chat("Show me users older than 30")
print(response)

# Reset conversation
agent.reset_conversation()
```

CLI Chat Interface

Start interactive chat:

```
iceframe-chat
```

Example session:

You: What tables do I have?

Assistant: You have the following tables:

- users
- orders
- products

You: Describe the users table

Assistant: The users table has the following schema:

- id (long, required)
- name (string)
- age (int)
- created_at (timestamp)

You: Show me users older than 30

Assistant: Here's a query to get users older than 30:

```
```python
from iceframe.expressions import Column

result = (ice.query("users")
 .filter(Column("age") > 30)
 .execute())
```

I found 15 users matching this criteria.

```
Commands

- `exit` or `quit`: End session
- `clear` or `reset`: Clear conversation history

Capabilities

Schema Discovery
```

You: What columns are in the orders table?

Assistant: The orders table contains:

order\_id (long)

user\_id (long)

amount (double)

status (string)

created\_at (timestamp)

```
Natural Language Queries
```

You: Find all orders over \$100

Assistant: [Executes query and shows results]

```
Code Generation
```

You: How do I join users with orders?

Assistant: Here's how to join users with orders:

```
result = (ice.query("users")
 .join("orders", on="user_id", how="inner")
 .select("name", "order_id", "amount")
 .execute())
```

```
Query Optimization
```

You: How can I make my queries faster?

Assistant: Based on your query patterns, I recommend:

Partition the orders table by date

Use column pruning to select only needed columns

Apply filters early to reduce data scanned

```
Supported LLMs
```

Provider	Models	Environment Variable
OpenAI	GPT-4, GPT-3.5-turbo	`OPENAI_API_KEY`
Anthropic	Claude 3.5 Sonnet, Claude 3 Opus	`ANTHROPIC_API_KEY`
Google	Gemini 2.0 Flash, Gemini Pro	`GOOGLE_API_KEY` or `GEMINI_API_KEY`

```
Extending with Custom LLMs
```

Create a custom LLM provider:

```
```python
from iceframe.agent.llm_base import BaseLLM, LLMConfig
```

```

class CustomLLM(BaseLLM):
    def chat(self, messages, tools=None):
        # Your implementation
        pass

    def stream_chat(self, messages):
        # Your implementation
        pass

# Use it
agent = IceFrameAgent(ice, llm=CustomLLM(config))

```

Source: `async.md`

Async Support

`IceFrame` provides `async` versions of core operations for non-blocking execution.

AsyncIceFrame

```

import asyncio
from iceframe.async_ops import AsyncIceFrame

async def main():
    config = {...}
    async_ice = AsyncIceFrame(config)

    # Async read
    df = await async_ice.read_table_async("users")

    # Async write
    await async_ice.append_to_table_async("users", new_data)

    # Async stats
    stats = await async_ice.stats_async("users")

asyncio.run(main())

```

Async Query Builder

```

from iceframe.expressions import Column

async def query_data():
    async_ice = AsyncIceFrame(config)

    query = await async_ice.query_async("users")
    result = await (query
        .filter(Column("age") > 25)

```

```
    .execute_async())

    return result

df = asyncio.run(query_data())
```

Use Cases

High Concurrency: Handle multiple table operations concurrently

Web Applications: Non-blocking API endpoints

Data Pipelines: Parallel processing of multiple tables

Source: branching.md

Branching and Tagging

IceFrame supports table branching and snapshot tagging (requires Pylceberg 0.6.0+ and catalog support).

Creating Branches

```
# Create branch from current snapshot
ice.create_branch("users", "experiment")

# Create branch from specific snapshot
ice.create_branch("users", "stable", snapshot_id=12345)
```

Tagging Snapshots

```
# Tag a snapshot for reference
table = ice.get_table("users")
snapshot_id = table.current_snapshot().snapshot_id

ice.tag_snapshot("users", snapshot_id, "v1.0")
```

Use Cases

Experimentation: Create branches for testing schema changes

Rollback: Tag stable snapshots for easy rollback

Versioning: Tag releases for reproducibility

Fast-Forwarding (Publishing)

You can fast-forward a branch (e.g., `main`) to another branch (e.g., `audit_branch`). This is essential for the Write-Audit-Publish (WAP) pattern.

```
from iceframe.branching import BranchManager

# Initialize manager
table = ice.get_table("users")
manager = BranchManager(table)

# Fast-forward main to audit_branch
manager.fast_forward("main", "audit_branch")
```

Write-Audit-Publish (WAP) Pattern

IceFrame supports the WAP pattern to ensure data quality:

Write: Write data to a branch (e.g., `audit_branch`).

Audit: Validate data in the branch.

Publish: Fast-forward `main` to the branch.

```
# 1. Write to branch
ice.append_to_table("users", new_data, branch="audit_branch")

# 2. Audit (Validate)
# ... run checks ...

# 3. Publish
manager.fast_forward("main", "audit_branch")
```

Source: catalog_ops.md

Catalog Operations

Manage catalog-level operations beyond basic table CRUD.

Registering Tables

You can register an existing Iceberg table (metadata.json) into the catalog. This is useful for:

Migrating tables between catalogs.

Recovering tables from storage.

Registering tables created by other engines.

```
# Register a table using its metadata location
metadata_url = "s3://bucket/warehouse/my_table/metadata/v1.metadata.json"

ice.register_table("my_new_table", metadata_url)
```

[!NOTE]

Not all catalogs support table registration. Check your catalog's documentation.

Source: cli.md

Command Line Interface (CLI)

IceFrame provides a CLI for managing Iceberg tables directly from the terminal.

Installation

The CLI requires the `cli` optional dependency:

```
pip install "iceframe[cli]"
```

Configuration

The CLI uses environment variables for configuration. You can set them in your shell or in a ` `.env` file.

```
export ICEBERG_CATALOG_URI="https://catalog.example.com"
export ICEBERG_CATALOG_TOKEN="your_token"
export ICEBERG_WAREHOUSE="s3://bucket/warehouse"
```

Commands

List Tables

List tables in a namespace (default is "default").

```
iceframe list
iceframe list --namespace marketing
```

Describe Table

Show table schema and partition spec.

```
iceframe describe my_table
```

Head Table

Show the first N rows of a table.

```
iceframe head my_table --n 10
```

Source: creating_tables.md

Creating Tables

IceFrame makes it easy to create Apache Iceberg tables with various schema formats.

Basic Usage

```
ice.create_table("my_table", schema)
```

Schema Formats

You can define schemas in several ways:

Dictionary Schema

Simple key-value pairs of column names and types.

```
schema = {  
    "id": "long",  
    "name": "string",  
    "price": "double",  
    "active": "boolean",  
    "created_at": "timestamp",  
    "birth_date": "date"  
}  
ice.create_table("products", schema)
```

Supported types: `string`, `int`, `long`, `float`, `double`, `boolean`, `timestamp`, `date`.

PyArrow Schema

For more control over types and nullability.

```
import pyarrow as pa  
  
schema = pa.schema([  
    pa.field("id", pa.int64(), nullable=False),  
    pa.field("name", pa.string()),
```

```
    pa.field("tags", pa.list_(pa.string())))
])
ice.create_table("users", schema)
```

Polars DataFrame

Infer schema from an existing DataFrame.

```
import polars as pl

df = pl.DataFrame({"id": [1], "name": ["test"]})
ice.create_table("inferred_table", df)
```

Namespaces

You can specify namespaces (databases/schemas) in the table name:

```
# Creates table 'sales' in 'marketing' namespace
ice.create_table("marketing.sales", schema)
```

If the namespace doesn't exist, IceFrame will attempt to create it.

Advanced Options

Partitioning

Partition data for better query performance.

```
# Not yet fully exposed in high-level API, use underlying PyIceberg table object
# or pass partition_spec to create_table (requires PyIceberg PartitionSpec object)
```

Table Properties

Set table properties like compression codec.

```
properties = {
    "write.parquet.compression-codec": "zstd"
}
ice.create_table("optimized_table", schema, properties=properties)
```

Source: data_quality.md

Data Quality

IceFrame includes a Data Validator to ensure data quality before or after operations.

Accessing Data Validator

```
validator = ice.validator
```

Checking for Nulls

Check if specific columns contain null values.

```
import polars as pl

df = pl.DataFrame(...)

if not ice.validator.check_nulls(df, ["id", "created_at"]):
    print("Data contains nulls in required columns!")
```

Validating Constraints

Validate data against SQL-like constraints or custom functions.

```
import polars as pl

df = pl.DataFrame(...)

results = ice.validator.validate(df, [
    pl.col("age") > 0,
    pl.col("status").is_in(["active", "inactive"])
])

if not results["passed"]:
    print("Validation failed:", results["details"])
```

Source: `data_quality_enhanced.md`

Enhanced Data Quality

IceFrame provides a robust suite of data validation methods to ensure your data meets quality standards before or after ingestion.

Usage

Access the validator via `ice.quality`. You can pass DataFrames, QueryBuilders, or SQL strings directly to expectation methods.

```
# 1. Using Query Builder
```

```

ice.quality.expect_column_values_to_be_unique(
    ice.query("my_table").filter("id > 100"),
    "id"
)

# 2. Using SQL (DataFusion) (requires DataFusion optional dependency)
ice.quality.expect_column_values_to_match_regex(
    "SELECT email FROM users WHERE active = true",
    "email",
    r"^.+@.+\.+\$"
)

# 3. Using Polars DataFrame directly
df = pl.DataFrame(...)
ice.quality.expect_column_values_to_not_be_null(df, "required_col")

```

Available Expectations

- `expect_column_values_to_be_unique(df, column)`
- `expect_column_values_to_be_between(df, column, min_value, max_value)`
- `expect_column_values_to_match_regex(df, column, regex)`
- `expect_column_values_to_be_in_set(df, column, value_set)`
- `expect_column_values_to_not_be_null(df, column)`
- `expect_table_row_count_to_be_between(df, min_value, max_value)`

Source: datafusion.md

DataFusion Integration

IceFrame integrates with Apache DataFusion to provide high-performance SQL execution on your Iceberg tables.

Installation

```
pip install "iceframe[datafusion]"
```

Usage

Use the `query_datafusion` method to execute SQL queries.

```

# Execute SQL query
df = ice.query_datafusion("SELECT * FROM my_table WHERE id > 100")

# Register multiple tables explicitly if needed (auto-registration is basic)

```

```
df = ice.query_datafusion(  
    "SELECT t1.id, t2.name FROM table1 t1 JOIN table2 t2 ON t1.id = t2.id",  
    tables=["table1", "table2"]  
)
```

Performance

DataFusion is an extensible query execution framework written in Rust that uses Apache Arrow as its in-memory format. It is often significantly faster than other engines for complex analytical queries, especially aggregations and joins, due to its vectorized execution engine.

Source: deleting_tables.md

Deleting Tables

IceFrame allows you to delete tables and remove specific rows.

Dropping Tables

Remove a table from the catalog.

```
ice.drop_table("temp_table")
```

This removes the table metadata from the catalog. The underlying data files may be retained depending on catalog configuration (GC properties).

Deleting Rows

Delete rows matching a filter expression.

```
# Delete all users from 'test' region  
ice.delete_from_table("users", "region = 'test'")
```

[!NOTE]

Row-level deletion requires Iceberg v2 tables and support from the underlying catalog and PyIceberg version.

Cleaning Up

After deleting tables or rows, you may want to perform [Maintenance](#) to clean up old files.

Source: dependencies.md

Dependencies

IceFrame is designed to be lightweight with a core set of dependencies and optional extras for specific features.

Core Dependencies

- `pyiceberg`: Core Iceberg client
- `polars`: High-performance DataFrame library
- `pyarrow`: Apache Arrow support
- `python-dotenv`: Environment variable management

Optional Dependencies

CLI (`[cli]`)

Required for the command-line interface.

- `typer`: CLI application builder
- `rich`: Terminal formatting

Install with:

```
pip install "iceframe[cli]"
```

Pydantic (`[pydantic]`)

Required for Pydantic integration.

- `pydantic>=2.0.0`: Data validation and settings management

Install with:

```
pip install "iceframe[pydantic]"
```

Ingestion (`[ingestion]`)

Required for various data ingestion and format support.

- `pdf`: PDF generation (`fpdf2`, `markdown-it-py`)
- `delta`: Delta Lake support (`deltalake`)
- `lance`: Lance support (`pylance`)
- `vortex`: Vortex support (`vortex-data`)

- ``excel``: Excel support (`fastexcel`)
- `gsheets`: Google Sheets support (`gspread`)
- `hudi`: Hudi support (`getdaft`)
- `sql`: SQL Database support (`connectorx`, `sqlalchemy`)
- `xml`: XML support (`lxml`)
- `stats`: Statistical file support (`pyreadstat`)
- `api`: REST API support (`requests`)
- `hf`: HuggingFace Datasets support (`datasets`)
- `html`: HTML table scraping (`lxml`, `html5lib`, `beautifulsoup4`)
- `clipboard`: Clipboard support (`pyperclip`)

Install with:

```
pip install "iceframe[ingestion]"
```

Or individually:

```
pip install "iceframe[sql]"
pip install "iceframe[xml]"
pip install "iceframe[stats]"
pip install "iceframe[api]"
pip install "iceframe[hf]"
pip install "iceframe[html]"
pip install "iceframe[clipboard]"
```

MCP Server (`[mcp]`)

Required for running the Model Context Protocol server.

- `mcp`: MCP server library

Install with:

```
pip install "iceframe[mcp]"
```

Notebook (`[notebook]`)

Required for Jupyter Notebook integration.

- `ipython>=8.0.0`: Interactive computing
- `ipywidgets>=8.0.0`: Interactive HTML widgets

Install with:

```
pip install "iceframe[notebook]"
```

Cloud Storage (`[aws]`, `[gcs]`, `[azure]`)

Required for accessing cloud storage backends.

`s3fs`: AWS S3 support

`gcsfs`: Google Cloud Storage support

`adlfs`: Azure Data Lake Storage support

Install with:

```
pip install "iceframe[aws]"
```

or

```
pip install "iceframe[aws,gcs]"
```

Async Support

Async operations use Python's built-in `asyncio` library (no additional dependencies required).

```
from iceframe.async_ops import AsyncIceFrame
```

AI Agent

The AI agent requires LLM provider packages. Install with:

```
pip install "iceframe[agent]"
```

This includes:

`openai>=1.0.0` - For GPT models

`anthropic>=0.18.0` - For Claude models

`google-generativeai>=0.3.0` - For Gemini models

`rich>=13.0.0` - For CLI formatting

You only need the API key for the LLM provider you want to use:

```
# Choose one:  
export OPENAI_API_KEY="your-key"  
export ANTHROPIC_API_KEY="your-key"  
export GOOGLE_API_KEY="your-key"
```

Scalability Features

```
pip install "iceframe[cache,streaming,monitoring]"
```

SQL Support (`[datafusion]`)

Required for high-performance SQL execution.

`datafusion>=35.0.0`: Apache DataFusion

```
pip install "iceframe[datafusion]"
```

Distributed Processing (`[distributed]`)

Required for distributed execution.

`ray>=2.0.0`: Ray

```
pip install "iceframe[distributed]"
```

Visualization (`[viz]`)

Required for generating charts.

`altair>=5.0.0`: Declarative statistical visualization

```
pip install "iceframe[viz]"
```

Source: distributed.md

Distributed Processing with Ray

IceFrame integrates with Ray to scale your data processing across a cluster of machines.

Installation

```
pip install "iceframe[distributed]"
```

Usage

Access the distributed executor via the `distribute` property.

```
# Initialize Ray (if not already running)
```

```
executor = ice.distribute

# Parallel Map
# Apply a function to a list of items in parallel
results = executor.map(
    lambda x: x * 2,
    [1, 2, 3, 4, 5]
)

# Parallel Table Reading
# Read multiple tables concurrently
dfs = executor.read_tables_parallel(
    ice_frame_config=ice.config,
    table_names=["table1", "table2", "table3"]
)
```

Configuration

The `RayExecutor` will automatically connect to an existing Ray cluster if available, or start a local one. You can pass arguments to `ray.init()` via the `RayExecutor` constructor if you instantiate it manually.

Source: [export.md](#)

Exporting Data

IceFrame allows you to export Iceberg table data to common file formats.

Export to Parquet

Parquet is efficient for analytics and storage.

```
ice.to_parquet("my_table", "output/data.parquet")
```

With filtering and column selection:

```
ice.to_parquet(
    "my_table",
    "output/us_sales.parquet",
    columns=["id", "amount"],
    filter_expr="region = 'US'"
)
```

Export to CSV

Useful for spreadsheets and simple data exchange.

```
ice.to_csv("my_table", "output/data.csv")
```

Export to JSON

Useful for web APIs and NoSQL databases.

```
ice.to_json("my_table", "output/data.json")
```

Performance Note

Exports work by reading the data into memory (as a Polars DataFrame) and then writing to disk. For extremely large tables, consider filtering the data first or using a distributed engine.

Source: incremental.md

Incremental Processing

IceFrame supports incremental reads and change data capture (CDC) to efficiently process only new or changed data.

Reading Incremental Data

Read only data added since a specific snapshot:

```
# Get current snapshot ID
table = ice.get_table("logs")
snapshot_id = table.current_snapshot().snapshot_id

# ... time passes, more data is added ...

# Read only new data
new_data = ice.read_incremental(
    "logs",
    since_snapshot_id=snapshot_id
)
```

Or use a timestamp:

```
import time

timestamp_ms = int(time.time() * 1000)
# ... later ...
new_data = ice.read_incremental(
    "logs",
    since_timestamp=timestamp_ms
)
```

Change Data Capture (CDC)

Track changes between two snapshots:

```
changes = ice.get_changes(  
    "users",  
    from_snapshot_id=snapshot1,  
    to_snapshot_id=snapshot2  
)  
  
print(f"Added: {changes['added'].height} rows")  
print(f"Deleted: {changes['deleted'].height} rows")
```

Source: ingest.md

Data Ingestion

IceFrame supports creating Iceberg tables directly from various data sources.

Installation

You can install the required dependencies for specific formats using optional extras:

```
# For Delta Lake  
pip install "iceframe[delta]"  
  
# For Lance  
pip install "iceframe[lance]"  
  
# For Vortex  
pip install "iceframe[vortex]"  
  
# For Excel  
pip install "iceframe[excel]"  
  
# For Google Sheets  
pip install "iceframe[gsheets]"  
  
# For Hudi  
pip install "iceframe[hudi]"
```

Usage

Delta Lake

Create an Iceberg table from a Delta Lake table.

```
ice.create_table_from_delta(
```

```
"my_namespace.from_delta",
"path/to/delta/table",
version=None # Optional version
)
```

Lance

Create an Iceberg table from a Lance dataset.

```
ice.create_table_from_lance(
    "my_namespace.from_lance",
    "path/to/lance/dataset"
)
```

Vortex

Create an Iceberg table from a Vortex file.

```
ice.create_table_from_vortex(
    "my_namespace.from_vortex",
    "path/to/vortex/file.vortex"
)
```

Excel

Create an Iceberg table from an Excel file.

```
ice.create_table_from_excel(
    "my_namespace.from_excel",
    "path/to/data.xlsx",
    sheet_name="Sheet1"
)
```

Google Sheets

Create an Iceberg table from a Google Sheet.

```
ice.create_table_from_gsheets(
    "my_namespace.from_gsheets",
    "https://docs.google.com/spreadsheets/d/...",
    credentials="path/to/credentials.json"
)
```

Apache Hudi

Create an Iceberg table from a Hudi table.

```
ice.create_table_from_hudi(  
    "my_namespace.from_hudi",  
    "path/to/hudi/table"  
)
```

Standard File Formats

IceFrame also supports standard file formats natively supported by Polars.

CSV

```
ice.create_table_from_csv(  
    "my_namespace.from_csv",  
    "path/to/data.csv",  
    has_header=True  
)
```

JSON

```
ice.create_table_from_json(  
    "my_namespace.from_json",  
    "path/to/data.json"  
)
```

Parquet

```
ice.create_table_from_parquet(  
    "my_namespace.from_parquet",  
    "path/to/data.parquet"  
)
```

IPC / Arrow

```
ice.create_table_from_ipc(  
    "my_namespace.from_ipc",  
    "path/to/data.arrow"  
)
```

Avro

```
ice.create_table_from_avro(  
    "my_namespace.from_avro",  
)
```

```
"path/to/data.avro"  
)
```

How it Works

These methods perform the following steps:

Read the data from the source into a Polars DataFrame.

Infer the schema from the DataFrame.

Create a new Iceberg table with that schema.

Append the data to the newly created table.

Source: ingest_advanced.md

Advanced File Ingestion

IceFrame supports advanced ingestion from SQL databases, XML files, and statistical software files (SAS, SPSS, Stata) through optional dependencies.

Supported Formats

Format	Extension	Install Command
---	---	---
SQL Database	(URI)	`pip install "iceframe[sql]"`
XML	`.xml`	`pip install "iceframe[xml]"`
SAS	`.sas7bdat`	`pip install "iceframe[stats]"`
SPSS	`.sav`	`pip install "iceframe[stats]"`
Stata	`.dta`	`pip install "iceframe[stats]"`

Usage

SQL Database

Read from any database supported by `connectorx` (Postgres, MySQL, SQLite, etc.) or `sqlalchemy`.

```
# Create table from SQL query  
ice.create_table_from_sql(  
    "my_table",  
    query="SELECT * FROM users WHERE active = true",  
    connection_uri="postgresql://user:pass@localhost:5432/db"  
)
```

XML

```
# Create table from XML
ice.create_table_from_xml("my_table", "data.xml")
```

Statistical Files (SAS, SPSS, Stata)

```
# Create table from SAS
ice.create_table_from_sas("my_table", "study_data.sas7bdat")

# Create table from SPSS
ice.create_table_from_spss("my_table", "survey.sav")

# Create table from Stata
ice.create_table_from_stata("my_table", "analysis.dta")
```

Generic Insert

You can also use `insert_from_file` with these formats.

```
ice.insert_from_file("my_table", "data.xml", format="xml")
ice.insert_from_file("my_table", "data.sas7bdat", format="sas")
```

Source: ingest_api.md

API Ingestion

IceFrame allows you to ingest data directly from REST APIs.

Usage

```
from iceframe import IceFrame
from iceframe.utils import load_catalog_config_from_env

config = load_catalog_config_from_env()
ice = IceFrame(config)

# Read from an API
url = "https://api.example.com/users"
ice.create_table_from_api("my_namespace.users", url)

# With a specific key for the list of records
ice.create_table_from_api(
    "my_namespace.users",
    "https://api.example.com/response",
    json_key="data"
)
```

```
# With authentication headers
ice.create_table_from_api(
    "my_namespace.secure_data",
    "https://api.example.com/secure",
    headers={"Authorization": "Bearer token"}
)
```

Dependencies

Requires `requests`.

```
pip install "iceframe[api]"
```

Source: ingest_clipboard.md

Clipboard Ingestion

IceFrame allows you to ingest data directly from your system clipboard. This is useful for ad-hoc analysis when copying data from Excel, Google Sheets, or other applications.

Usage

Copy data to your clipboard (e.g., select cells in Excel and press Ctrl+C).

Run the following code:

```
from iceframe import IceFrame
from iceframe.utils import load_catalog_config_from_env

config = load_catalog_config_from_env()
ice = IceFrame(config)

# Create table from clipboard content
ice.create_table_from_clipboard("my_namespace.clipboard_data")
```

Dependencies

Requires `pyperclip` (and `pandas`).

```
pip install "iceframe[clipboard]"
```

Source: ingest_folder.md

Folder Ingestion

IceFrame allows you to ingest multiple files from a folder into a single table.

Usage

```
from iceframe import IceFrame
from iceframe.utils import load_catalog_config_from_env

config = load_catalog_config_from_env()
ice = IceFrame(config)

# Read all CSVs from a folder
ice.create_table_from_folder(
    "my_namespace.daily_logs",
    "/path/to/logs",
    pattern="*.csv"
)

# Read all Parquet files
ice.create_table_from_folder(
    "my_namespace.archive",
    "/path/to/archive",
    pattern="*.parquet"
)
```

Supported Formats

CSV

JSON

Parquet

Excel

The format is inferred from the file extension.

Source: ingest_html.md

HTML Ingestion

IceFrame allows you to scrape tables from HTML pages.

Usage

```
from iceframe import IceFrame
from iceframe.utils import load_catalog_config_from_env

config = load_catalog_config_from_env()
ice = IceFrame(config)
```

```
# Read tables from a URL
ice.create_table_from_html(
    "my_namespace.wikipedia_data",
    "https://en.wikipedia.org/wiki/List_of_countries_by_GDP_(nominal)"
)

# Match a specific table
ice.create_table_from_html(
    "my_namespace.specific_table",
    "https://example.com/data",
    match="Population"
)
```

Dependencies

Requires `lxml`, `html5lib`, and `beautifulsoup4`.

```
pip install "iceframe[html]"
```

Source: ingest_huggingface.md

HuggingFace Ingestion

IceFrame allows you to ingest datasets directly from the HuggingFace Hub.

Usage

```
from iceframe import IceFrame
from iceframe.utils import load_catalog_config_from_env

config = load_catalog_config_from_env()
ice = IceFrame(config)

# Read a dataset
ice.create_table_from_huggingface(
    "my_namespace.imdb",
    "imdb",
    split="train"
)

# With specific configuration
ice.create_table_from_huggingface(
    "my_namespace.glue_mrpc",
    "glue",
    name="mrpc",
    split="validation"
)
```

Dependencies

Requires `datasets`.

```
pip install "iceframe[hf]"
```

Source: ingest_native.md

Native File Ingestion

IceFrame supports ingesting data from various file formats that are natively supported by the underlying engines (Polars, PyArrow) without requiring additional dependencies.

Supported Formats

CSV (`.csv`)

JSON (`.json`, `ndjson`)

Parquet (`.parquet`)

IPC / Arrow / Feather (`.ipc`, `arrow`, `feather`)

Avro (`.avro`)

ORC (`.orc`)

Usage

Creating a Table from a File

You can create a new Iceberg table directly from a file. The schema is inferred from the file content.

```
# Create from Parquet
ice.create_table_from_parquet("my_namespace.table_from_parquet", "data.parquet")

# Create from CSV
ice.create_table_from_csv("my_namespace.table_from_csv", "data.csv")

# Create from JSON
ice.create_table_from_json("my_namespace.table_from_json", "data.json")

# Create from ORC
ice.create_table_from_orc("my_namespace.table_from_orc", "data.orc")
```

Inserting Data from a File

You can insert data from a file into an existing table using the `insert_from_file` method. This method automatically detects the file format based on the extension, or you can specify it explicitly.

```
# Insert from CSV (format inferred)
ice.insert_from_file("my_table", "new_data.csv")

# Insert from JSON with explicit format
ice.insert_from_file("my_table", "new_data.json", format="json")

# Insert into a specific branch
ice.insert_from_file("my_table", "experiment_data.parquet", branch="experiment")
```

Supported Arguments

All ingestion methods accept **kwargs** which are passed directly to the underlying Polars read functions (e.g., `pl.read_csv`, `pl.read_parquet`). Refer to the [Polars documentation for available options](#).

```
# Read CSV with specific options
ice.create_table_from_csv(
    "my_table",
    "data.csv",
    has_header=True,
    separator=";"
)
```

Source: ingest_optional.md

Optional File Ingestion

IceFrame supports additional file formats through optional dependencies. These formats require installing specific extras.

Supported Formats

Format	Extension	Install Command
---	---	---
Excel	`.xlsx`, `.xls`	`pip install "iceframe[excel]"`
Delta Lake	(Directory)	`pip install "iceframe[delta]"`
Lance	`.lance`	`pip install "iceframe[lance]"`
Vortex	`.vortex`	`pip install "iceframe[vortex]"`
Google Sheets	(URL)	`pip install "iceframe[gsheets]"`
Hudi	(Directory)	`pip install "iceframe[hudi]"`

Usage

Excel

```
# Create table from Excel
ice.create_table_from_excel(
    "my_table",
    "data.xlsx",
    sheet_name="Sales_2024"
)
```

Delta Lake

```
# Create table from Delta Lake
ice.create_table_from_delta(
    "my_table",
    "/path/to/delta/table",
    version=1
)
```

Google Sheets

```
# Create table from Google Sheets
ice.create_table_from_gsheets(
    "my_table",
    "https://docs.google.com/spreadsheets/d/...",
    credentials="path/to/service_account.json"
)
```

Generic Insert

You can also use `insert_from_file` with these formats, provided the dependencies are installed.

```
ice.insert_from_file("my_table", "data.xlsx", format="excel")
ice.insert_from_file("my_table", "/path/to/delta", format="delta")
```

Source: ingestion.md

Data Ingestion & Bulk Import

IceFrame supports efficient bulk import of existing data files into Iceberg tables without rewriting them.

Adding Files

If you have existing Parquet, Avro, or ORC files (e.g., from a migration or another job), you can register them directly into the table. This is a metadata-only operation and is extremely fast.

```
files = [
    "s3://bucket/data/file1.parquet",
    "s3://bucket/data/file2.parquet"
]

# Add files to the table
ice.add_files("my_table", files)
```

Requirements

Files must match the table's schema.

Files must be in a location accessible by the catalog/engine.

If the table is partitioned, files should ideally align with partitions, though Iceberg can handle unpartitioned files in partitioned tables (they will be scanned to determine partition values if metrics are available).

[!TIP]

Use this for migrating large datasets from legacy systems or other table formats.

Source: joins.md

JOIN Support

IceFrame Query Builder supports cross-table joins.

Basic JOIN

```
from iceframe.expressions import Column

# Inner join
result = (ice.query("users")
    .join("orders", on="user_id", how="inner")
    .select("name", "order_id", "amount")
    .execute())
```

JOIN Types

Supported join types: `inner`, `left`, `right`, `outer`

```
# Left join
result = (ice.query("users")
    .join("orders", on="user_id", how="left")
    .execute())
```

```
# Multiple joins
result = (ice.query("users")
    .join("orders", on="user_id")
    .join("products", on="product_id")
    .select("name", "product_name", "amount")
    .execute())
```

JOIN with Filters

```
result = (ice.query("users")
    .join("orders", on="user_id")
    .filter(Column("amount") > 100)
    .select("name", "order_id")
    .execute())
```

Source: lazy_reading.md

Lazy Reading

IceFrame supports true lazy reading of Iceberg tables, allowing you to process datasets larger than memory by streaming data in batches.

Usage

Use the `read_table_chunked` method from the memory manager (or via `ice.memory`) to iterate over the table in chunks.

```
from iceframe.memory import MemoryManager

# Initialize memory manager with a limit (optional)
mem = MemoryManager(max_memory_mb=1024) # 1GB limit

# Iterate over the table
for chunk in mem.read_table_chunked(ice, "my_huge_table"):
    # Process chunk (Polars DataFrame)
    print(chunk.head())

    # Perform aggregations, write to another table, etc.
```

How it Works

Unlike standard reading which loads the entire table into memory, lazy reading uses Pylceberg's batch reader to stream Arrow batches from storage. These batches are converted to Polars DataFrames on the fly, ensuring that only a small portion of the data is in memory at any given time.

Memory Safety

The `MemoryManager` can be configured with a `max_memory_mb` limit. It checks memory usage before yielding each chunk and raises a `MemoryError` if the limit is exceeded, helping prevent OOM crashes.

Source: maintenance.md

Table Maintenance

Iceberg tables require periodic maintenance to ensure optimal performance and manage storage costs. IceFrame provides simple methods for common maintenance tasks.

Expiring Snapshots

Remove old table snapshots to free up space and keep metadata size manageable.

```
# Remove snapshots older than 7 days, keeping at least the last 1
ice.expire_snapshots("my_table", older_than_days=7, retain_last=1)
```

Removing Orphan Files

Clean up data files that are no longer referenced by any snapshot (e.g., from failed writes).

```
# Remove orphan files older than 3 days
ice.remove_orphan_files("my_table", older_than_days=3)
```

Compacting Data Files

Combine small data files into larger ones to improve read performance (compaction).

```
# Compact files to target size of 512 MB
ice.compact_data_files("my_table", target_file_size_mb=512)
```

[!NOTE]

Compaction Implementation Details

IceFrame currently uses a manual compaction strategy (reading data, optionally sorting, and rewriting).

This process typically generates two distinct snapshots in the table history:

A **Delete** snapshot (removing old files)

An **Append** snapshot (adding new compacted files)

This is expected behavior until upstream libraries support atomic `rewrite_data_files`

operations natively. Functionally, the data is correctly compacted and replaced.

[!TIP]

Run compaction regularly on tables with frequent small updates (streaming ingestion).

Best Practices

Schedule Maintenance: Run these operations periodically (e.g., daily or weekly) via a scheduler like Airflow.

Order of Operations:

- `expire_snapshots`
- `remove_orphan_files`
- `compact_data_files`

Source: [mcp.md](#)

MCP Server Integration

IceFrame includes a Model Context Protocol (MCP) server that exposes its capabilities to AI assistants and IDEs (like Claude Desktop, Cursor, etc.) via stdio.

Installation

Install IceFrame with the `mcp` extra:

```
pip install "iceframe[mcp]"
```

Configuration

The MCP server requires the same environment variables as the IceFrame library to connect to your Iceberg catalog.

- `ICEBERG_CATALOG_URI` (Required)
- `ICEBERG_CATALOG_TYPE` (Default: `rest`)
- `ICEBERG_WAREHOUSE`
- `ICEBERG_TOKEN`
- `ICEBERG_CREDENTIAL`
- `ICEBERG_OAUTH2_SERVER_URI`

Usage

Getting Configuration

To get the JSON configuration for your MCP client (e.g., for `claude_desktop_config.json`), run:

```
iceframe mcp config
```

This will output a JSON object like:

```
{  
  "mcpServers": {  
    "iceframe": {  
      "command": "/path/to/python",  
      "args": [  
        "-m",  
        "iceframe.cli",  
        "mcp",  
        "start"  
      ],  
      "env": {  
        "ICEBERG_CATALOG_URI": "..."  
      }  
    }  
  }  
}
```

Copy this configuration into your client's settings file.

Starting the Server

The server is typically started automatically by the MCP client using the command specified in the configuration. However, you can start it manually for testing:

```
iceframe mcp start
```

Available Tools

The MCP server exposes the following tools to the AI assistant:

- `list_tables(namespace)` : List tables in a namespace.
- `describe_table(table_name)` : Get schema and metadata for a table.
- `get_table_stats(table_name)` : Get table statistics.
- `execute_query(table_name, query, limit)` : Execute a query (filter expression) on a table.
- `generate_code(operation)` : Generate Python code for complex operations.
- `generate_sql(description)` : Generate SQL query templates.

`list_documentation()` : List available documentation files.

`read_documentation(page)` : Read the content of a documentation file.

Source: namespaces.md

Namespace Management

IceFrame allows you to manage namespaces (schemas/databases) in your Iceberg catalog.

Accessing Namespace Manager

```
# Access via the namespaces property
ns_manager = ice.namespaces
```

Creating a Namespace

```
ice.create_namespace("marketing", {"owner": "team-marketing"})
```

Dropping a Namespace

```
ice.drop_namespace("marketing")
```

Listing Namespaces

```
# List top-level namespaces
namespaces = ice.list_namespaces()
print(namespaces)

# List nested namespaces
sub_namespaces = ice.list_namespaces("marketing")
```

Source: native_maintenance.md

Native Maintenance Operations

IceFrame provides native implementations of critical maintenance operations that work independently of Pylceberg version or catalog support.

Native Orphan File Removal

Remove data files that are no longer referenced by any snapshot.

```
from iceframe.gc import GarbageCollector

# Get table
table = ice.get_table("my_table")
gc = GarbageCollector(table)

# Dry run - list orphans without deleting
orphans = gc.remove_orphan_files(dry_run=True)
print(f"Found {len(orphans)} orphaned files")

# Remove orphans older than 3 days
import time
three_days_ago_ms = int((time.time() - 3 * 86400) * 1000)
removed = gc.remove_orphan_files(older_than_ms=three_days_ago_ms)
print(f"Removed {len(removed)} orphaned files")
```

How It Works

Scan Manifests: Reads all manifest files from the current snapshot to build a set of referenced data files

List Storage: Lists all files in the table's data directory

Find Orphans: Computes the difference (files in storage - files in manifests)

Age Filter: Optionally filters by file modification time

Delete: Removes orphaned files (unless `dry_run=True`)

Use Cases

After Failed Writes: Clean up files from failed write operations

After Compaction: Remove old data files after rewriting

Storage Optimization: Reclaim storage from deleted/replaced files

[!TIP]

Always run with `dry_run=True` first to verify which files will be removed.

Native Snapshot Expiration

Expire old snapshots to reduce metadata size and improve query planning performance.

```
from iceframe.gc import GarbageCollector
import time

table = ice.get_table("my_table")
gc = GarbageCollector(table)
```

```

# Expire snapshots older than 7 days, keeping at least 5
seven_days_ago_ms = int((time.time() - 7 * 86400) * 1000)

try:
    expired = gc.expire_snapshots(
        older_than_ms=seven_days_ago_ms,
        retain_last=5
    )
    print(f"Expired {len(expired)} snapshots")
except NotImplementedError as e:
    print(f"Snapshot expiration not supported: {e}")

```

How It Works

List Snapshots: Gets all snapshots from table metadata

Apply Retention: Filters snapshots based on age and retention count

Expire: Uses Pylceberg's `expire_snapshots` if available, otherwise raises `NotImplementedError`

[!NOTE]

Snapshot expiration requires Pylceberg 0.7.0+ or catalog support. If not available, the operation will raise `NotImplementedError`.

Use Cases

Metadata Optimization: Reduce table metadata size

Query Performance: Improve query planning by reducing snapshot count

Compliance: Meet data retention policies

Comparison with Pylceberg

Operation Pylceberg IceFrame Native		
:----- :----- :-----		
Orphan File Removal `table.remove_orphan_files()` (v0.7+) Works on any version		
Snapshot Expiration `table.expire_snapshots()` (catalog-dependent) ▲ Wraps Pylceberg, adds retry logic		
Dry Run Support Limited Full support		
Age Filtering Basic Enhanced with file stat checks		

Best Practices

1. Schedule Regular Maintenance

```

# Weekly orphan file cleanup
gc.remove_orphan_files(older_than_ms=three_days_ago_ms)

```

```
# Monthly snapshot expiration
gc.expire_snapshots(older_than_ms=thirty_days_ago_ms, retain_last=10)
```

2. Use Dry Run First

```
# Always verify before deleting
orphans = gc.remove_orphan_files(dry_run=True)
if len(orphans) > 1000:
    print("Warning: Large number of orphans detected!")
    # Investigate before proceeding
```

3. Monitor Storage Savings

```
import os

# Calculate storage before
orphans = gc.remove_orphan_files(dry_run=True)
total_size = sum(os.path.getsize(f) for f in orphans if os.path.exists(f))
print(f"Potential savings: {total_size / 1024**3:.2f} GB")

# Perform cleanup
gc.remove_orphan_files()
```

Limitations

Snapshot Expiration: Requires Pylceberg 0.7.0+ or catalog support

Concurrent Writes: Orphan detection may miss files from concurrent write operations

Distributed Storage: File listing performance depends on storage system (S3, HDFS, etc.)

See Also

[Maintenance Guide](#)

[Garbage Collection](#)

[Table Optimization](#)

Source: [notebooks.md](#)

Notebook Integration

IceFrame provides rich integration with Jupyter Notebooks and IPython environments.

Installation

To use notebook features, install IceFrame with the `notebook` extra:

```
pip install "iceframe[notebook]"
```

Rich Display

When you display an `IceFrame` instance in a notebook, it shows a formatted summary of the connection and available namespaces.

```
from iceframe import IceFrame  
  
ice = IceFrame(config)  
ice # Displays HTML summary
```

Magic Commands

IceFrame includes IPython magic commands to simplify interaction.

Loading the Extension

First, load the extension:

```
%load_ext iceframe.magics
```

%iceframe

Set the active IceFrame instance for magic commands.

```
ice = IceFrame(config)  
%iceframe ice
```

Check status:

```
%iceframe status
```

%%iceql

Execute SQL queries directly in a cell using the active IceFrame instance.

```
%%iceql  
SELECT * FROM my_table LIMIT 10
```

You can also perform joins:

```
%%iceql
SELECT
    t1.name,
    t2.order_total
FROM users t1
JOIN orders t2 ON t1.id = t2.user_id
```

Note: `%%iceql` uses Polars SQL context under the hood. It automatically registers tables referenced in the query from your Iceberg catalog.

Source: partitioning.md

Partition Management

IceFrame allows you to manage table partitioning to optimize query performance.

Accessing Partition Manager

```
# Get partition manager for a table
partitioner = ice.partition_by("logs")
```

Adding Partition Fields

IceFrame supports various transforms: `identity`, `bucket`, `truncate`, `year`, `month`, `day`, `hour`.

```
# Partition by 'category' (identity)
ice.partition_by("logs").add_partition_field("category")

# Partition by day of 'timestamp'
ice.partition_by("logs").add_partition_field("timestamp", "day", name="day_ts")

# Partition by bucket of 'user_id'
ice.partition_by("logs").add_partition_field("user_id", "bucket", name="user_bucket")
```

16,

Dropping Partition Fields

```
# Drop partition field
ice.partition_by("logs").drop_partition_field("category")
```

Source: pydantic.md

Pydantic Integration

IceFrame supports integration with [Pydantic](#) (v2) for schema definition and data validation.

Installation

To use Pydantic features, install IceFrame with the `pydantic` extra:

```
pip install "iceframe[pydantic]"
```

Defining Tables with Pydantic Models

You can use Pydantic models to define your table schema instead of PyArrow or Pylceberg schemas.

```
from pydantic import BaseModel
from typing import Optional
from datetime import datetime
from iceframe import IceFrame

# Define your model
class User(BaseModel):
    id: int
    name: str
    email: Optional[str] = None
    created_at: datetime = datetime.now()
    is_active: bool = True

# Initialize IceFrame
ice = IceFrame(config)

# Create table using the model
ice.create_table("my_namespace.users", schema=User)
```

Inserting Data

You can insert a list of Pydantic model instances directly into a table using `insert_items`.

```
# Create user instances
users = [
    User(id=1, name="Alice", email="alice@example.com"),
    User(id=2, name="Bob", is_active=False)
]

# Insert into table
ice.insert_items("my_namespace.users", users)
```

Type Mapping

IceFrame maps Python/Pydantic types to Iceberg types as follows:

Python Type		Iceberg Type
`---		---
`str`		`string`
`int`		`long`
`float`		`double`
`bool`		`boolean`
`datetime`		`timestamp`
`date`		`date`
`List[T]`		`list<T>`
`Optional[T]`		`T` (nullable)

Advanced Usage

For more complex schemas (nested structs, maps), you can nest Pydantic models.

```
class Address(BaseModel):
    street: str
    city: str
    zip: str

class UserWithAddress(BaseModel):
    id: int
    name: str
    address: Address # Maps to Iceberg Struct
```

Source: query_builder.md

Query Builder API

IceFrame provides a powerful, fluent Query Builder API for constructing complex queries with SQL-like capabilities.

Overview

The Query Builder allows you to:

Select columns and apply expressions

Filter data with predicate pushdown support

Group by and aggregate data

Sort and limit results

Use window functions and case statements

Perform write operations (Insert, Update, Delete, Merge)

Basic Usage

Start a query using `ice.query("table_name")`:

```
from iceframe.expressions import col, lit

df = (ice.query("sales")
      .select("id", "amount")
      .filter(col("amount") > 100)
      .execute())
```

Expressions

IceFrame provides a unified expression system that works with both PyIceberg (for pushdown) and Polars (for local processing).

```
from iceframe.expressions import col, lit

# Binary operations
col("age") > 18
col("status") == "active"

# Boolean logic
(col("age") > 18) & (col("status") == "active")
(col("category") == "A") | (col("category") == "B")

# IN / IS NULL
col("id").is_in([1, 2, 3])
col("name").is_null()
```

Aggregations

Use standard SQL aggregate functions:

```
from iceframe.functions import count, sum, avg, min, max

df = (ice.query("sales")
      .select(
          col("region"),
          sum(col("amount")).alias("total_sales"),
          avg(col("amount")).alias("avg_sales")
      )
      .group_by("region")
      .execute())
```

Window Functions

Support for window functions like `row_number`, `rank`, `dense_rank`:

```
from iceframe.functions import row_number

df = (ice.query("sales")
      .select(
          col("id"),
          col("amount"),
          row_number().over(
              partition_by=col("region"),
              order_by=col("amount"))
      ).alias("rank")
)
.execute()
```

Case Statements

Conditional logic with `when/otherwise`:

```
from iceframe.functions import when

df = (ice.query("users")
      .select(
          col("name"),
          when(col("age") < 18, "Minor")
            .when(col("age") < 65, "Adult")
            .otherwise("Senior")
      ).alias("age_group")
)
.execute()
```

Write Operations

The Query Builder also supports write operations.

Insert

```
ice.query("users").insert(new_data_df)
```

Update

Update rows matching a filter:

```
(ice.query("users")
  .filter(col("id") == 123)
  .update({"status": "inactive"}))
```

[!WARNING]

Updates are currently implemented as Copy-on-Write (overwrite entire table). Use with caution on large tables.

Delete

Delete rows matching a filter:

```
(ice.query("users")
  .filter(col("status") == "deleted")
  .delete())
```

Merge (Upsert)

Merge source data into the target table:

```
(ice.query("target_table")
  .merge(
    source_data=source_df,
    on="id",
    when_matched_update={"status": "status", "updated_at": "updated_at"},
    when_not_matched_insert={"id": "id", "status": "status", "created_at": "created_at"})
  )
```

Source: reading_tables.md

Reading Tables

IceFrame provides a simple API to read Iceberg tables into Polars DataFrames.

Basic Reading

```
df = ice.read_table("my_table")
```

Column Selection

Read only specific columns to improve performance.

```
df = ice.read_table("users", columns=["id", "email"])
```

Filtering

Filter data at the source (predicate pushdown).

```
# Filter expression using SQL-like syntax
```

```
df = ice.read_table("sales", filter_expr="amount > 100 AND region = 'US'")
```

Limiting Results

Limit the number of rows returned.

```
df = ice.read_table("logs", limit=100)
```

Time Travel

Read the table as it existed at a specific point in time.

By Snapshot ID

```
df = ice.read_table("my_table", snapshot_id=123456789012345)
```

By Timestamp

```
# Read as of 1 hour ago
timestamp_ms = int((time.time() - 3600) * 1000)
df = ice.read_table("my_table", as_of_timestamp=timestamp_ms)
```

Accessing Underlying Table

For advanced operations, you can access the underlying PyIceberg Table object.

```
table = ice.get_table("my_table")
# Use PyIceberg API directly
scan = table.scan()
```

Source: `data_quality_gate.md`

Data Quality Gate Recipe

This recipe demonstrates how to implement a "Write-Audit-Publish" pattern using IceFrame's branching and data quality features.

Scenario

You want to load data into a production table, but only if it passes strict quality checks (e.g., no nulls in critical columns, values within range). If checks fail, the data should be

isolated for review without affecting production.

Implementation

```
from iceframe import IceFrame
from iceframe.branching import BranchManager

ice = IceFrame(config)
table_name = "finance.transactions"
staging_branch = "staging_audit"

def load_with_quality_gate(new_data_df):
    # 1. Create a Branch for Staging
    # We write to a branch first so production readers don't see incomplete/bad data
    branch_manager = BranchManager(ice.catalog)

    # Create or replace branch pointing to current main
    try:
        branch_manager.create_branch(table_name, staging_branch)
    except:
        # If exists, fast-forward or reset (simplified here)
        pass

    print(f"Writing data to branch '{staging_branch}'...")

    # 2. Write to the Branch
    # (IceFrame write methods support a 'branch' argument if implemented,
    # or we configure the write to target the branch ref)
    # For this recipe, we assume we can write to the branch:
    ice.append_to_table(table_name, new_data_df, branch=staging_branch)

    # 3. Run Data Quality Checks on the Branch
    print("Running quality checks...")

    # Define constraints
    constraints = [
        {"type": "not_null", "columns": ["transaction_id", "amount"]},
        {"type": "range", "column": "amount", "min": 0},
        {"type": "unique", "column": "transaction_id"}
    ]

    # Validate data in the branch
    # (We read from the branch to validate)
    validation_results = ice.validate_data(
        table_name,
        constraints,
        branch=staging_branch
    )

    if validation_results["passed"]:
        print(" Quality checks passed. Promoting to main.")

    # 4. Fast-Forward 'main' to 'staging_branch'
    # This makes the data visible to production users atomically
```

```

branch_manager.fast_forward(table_name, "main", staging_branch)

else:
    print(" Quality checks FAILED.")
    print(f"Violations: {validation_results['violations']}") 
    print(f"Bad data is isolated in branch '{staging_branch}' for review.") 
    # Do NOT merge to main. Alert the team.

# Example Run
df = pl.read_parquet("incoming_transactions.parquet")
load_with_quality_gate(df)

```

Key Features Used

Branching: Isolates unverified data from production readers.

Data Validator: Automates quality checks.

Atomic Promotion: Fast-forward merge ensures all-or-nothing visibility.

Source: etl_pipeline.md

ETL Pipeline Recipe

This recipe demonstrates how to build a simple ETL (Extract, Transform, Load) pipeline using IceFrame.

Scenario

You need to ingest raw JSON logs, clean the data, enrich it, and load it into an Iceberg table partitioned by date.

Implementation

```

from iceframe import IceFrame
import polars as pl
from datetime import datetime

# 1. Initialize IceFrame
config = {
    "uri": "http://localhost:8181",
    "type": "rest",
    "warehouse": "s3://warehouse"
}
ice = IceFrame(config)

def run_etl_job(source_file: str, target_table: str):
    print(f"Starting ETL job for {source_file}...")

```

```

# --- EXTRACT ---
# Read raw data using Polars
raw_df = pl.read_json(source_file)

# --- TRANSFORM ---
# Clean and enrich data
processed_df = (
    raw_df
    .with_columns([
        # Convert timestamp string to datetime
        pl.col("timestamp").str.to_datetime(),
        # Extract date for partitioning
        pl.col("timestamp").str.to_datetime().dt.date().alias("event_date"),
        # Clean strings
        pl.col("user_agent").str.strip_chars(),
        # Add processing metadata
        pl.lit(datetime.now()).alias("processed_at")
    ])
    .filter(
        # Remove invalid records
        pl.col("user_id").is_not_null()
    )
)

# --- LOAD ---
# Check if table exists, create if not
if not ice.table_exists(target_table):
    print(f"Creating table {target_table}...")
    ice.create_table(
        target_table,
        schema=processed_df.schema,
        partition_spec=[("event_date", "identity")]
    )

# Append data to Iceberg table
print(f"Writing {processed_df.height} records to {target_table}...")
ice.append_to_table(target_table, processed_df)

print("ETL job completed successfully!")

# Run the job
run_etl_job("raw_logs.json", "analytics.web_logs")

```

Key Features Used

Polars Integration: Uses Polars for efficient in-memory transformation.

Schema Inference: Automatically infers Iceberg schema from the DataFrame.

Partitioning: Creates a partitioned table for optimized querying.

Idempotency: Checks for table existence before creation.

Source: incremental_ingestion.md

Incremental Ingestion Recipe

This recipe demonstrates how to process data incrementally, reading only what has changed since the last run.

Scenario

You have a downstream table `daily_summary` that aggregates data from an upstream `raw_events` table. You want to run this aggregation periodically, processing only new data added to `raw_events`.

Implementation

```
from iceframe import IceFrame
import json
import os

ice = IceFrame(config)
state_file = "ingestion_state.json"

def get_last_snapshot_id():
    if os.path.exists(state_file):
        with open(state_file, "r") as f:
            return json.load(f).get("last_snapshot_id")
    return None

def save_state(snapshot_id):
    with open(state_file, "w") as f:
        json.dump({"last_snapshot_id": snapshot_id}, f)

def run_incremental_job():
    source_table = "raw_events"
    target_table = "daily_summary"

    # 1. Get last processed state
    last_snapshot = get_last_snapshot_id()

    # 2. Read incremental data
    print(f"Reading {source_table} since snapshot {last_snapshot}...")

    # Use IceFrame's incremental reader
    # This returns a DataFrame containing only new rows appended since snapshot
    new_data = ice.read_incremental(
        source_table,
        since_snapshot_id=last_snapshot
    )

    if new_data.height == 0:
        print("No new data found.")
```

```

    return

# 3. Process the data (Aggregation)
summary_df = (
    new_data
    .group_by("event_date", "category")
    .agg([
        pl.count().alias("event_count"),
        pl.sum("amount").alias("total_amount")
    ])
)

# 4. Write to target (Upsert/Merge)
# We merge into the summary table to update counts
(ice.query(target_table)
    .merge(summary_df, on=["event_date", "category"])
    .when_matched_update({
        "event_count": pl.col("target.event_count") + pl.col("source.event_count"),
        "total_amount": pl.col("target.total_amount") +
pl.col("source.total_amount")
    })
    .when_not_matched_insert({
        "event_date": pl.col("source.event_date"),
        "category": pl.col("source.category"),
        "event_count": pl.col("source.event_count"),
        "total_amount": pl.col("source.total_amount")
    })
    .execute())

# 5. Update state
# Get the current snapshot ID of the source table to save for next time
current_snapshot = ice.get_table(source_table).current_snapshot().snapshot_id
save_state(current_snapshot)
print(f"Job finished. State updated to snapshot {current_snapshot}")

run_incremental_job()

```

Key Features Used

Incremental Read: `read_incremental` efficiently fetches only new data.

Merge/Upsert: Updates existing aggregates or inserts new ones.

State Management: Tracks progress via snapshot IDs.

Source: scd_type_2.md

Slowly Changing Dimensions (SCD) Type 2 Recipe

This recipe demonstrates how to implement SCD Type 2 (retaining full history) using IceFrame's merge capabilities.

Scenario

You have a `users` table and receive updates. You want to track historical changes to user profiles (e.g., address changes) by creating new records for updates while marking old records as inactive.

Implementation

```
from iceframe import IceFrame
from iceframe.expressions import Column
import polars as pl

ice = IceFrame(config)
target_table = "dim.users"

def process_scd_type_2(updates_df: pl.DataFrame):
    """
    Apply SCD Type 2 updates to the users dimension table.
    Schema: user_id, name, address, is_active, valid_from, valid_to
    """

    # 1. Identify records that have actually changed
    # Read current active records
    current_df = (
        ice.query(target_table)
        .filter(Column("is_active") == True)
        .execute()
    )

    # Join to find changes
    # (In a real scenario, you'd hash columns to detect changes efficiently)

    # 2. Prepare the MERGE operation
    # For SCD Type 2, we typically do this in two steps or use a complex merge.
    # Here is a simplified approach using IceFrame's merge:

    # Step A: Expire old records
    # Update existing records where ID matches but content differs
    (ice.query(target_table)
        .merge(updates_df, on="user_id")
        .when_matched_update({
            "is_active": False,
            "valid_to": pl.col("source.valid_from")
        })
        .execute())

    # Step B: Insert new versions
    # Insert new records for all updates
    new_records = updates_df.with_columns([
        pl.lit(True).alias("is_active"),
```

```

    pl.lit(None).alias("valid_to")
])

ice.append_to_table(target_table, new_records)

# Example Usage
updates = pl.DataFrame({
    "user_id": [101, 102],
    "name": ["Alice", "Bob"],
    "address": ["New Address St", "Same Address"],
    "valid_from": [datetime.now(), datetime.now()]
})

process_scd_type_2(updates)

```

[!NOTE]

True SCD Type 2 often requires complex logic to handle out-of-order data and exact timestamp alignment. This recipe shows the fundamental pattern of expiring old rows and inserting new ones.

Key Features Used

Merge Operation: Used to update existing records based on a key.

Predicate Pushdown: Efficiently reads only active records.

Source: rollback.md

Rollback & Snapshot Management

IceFrame provides tools to manage table history, allowing you to rollback to previous states or manage branch pointers.

Rollback

Revert the table state to a specific snapshot or point in time.

```

# Rollback to a specific snapshot ID
ice.rollback_to_snapshot("my_table", 1234567890)

# Rollback to a timestamp (milliseconds since epoch)
ice.rollback_to_timestamp("my_table", 1704067200000)

```

Snapshot Management

Explicitly set the current snapshot of the table. This is useful for advanced workflows like cherry-picking or manually moving branch pointers.

```
# Set the current snapshot of the table
ice.call_procedure("my_table", "set_current_snapshot", snapshot_id=9876543210)
```

[!WARNING]

Rolling back changes the current state of the table. Ensure you have the correct snapshot ID or timestamp before proceeding.

Source: scalability.md

Scalability Features

IceFrame includes comprehensive scalability features for high-performance data processing.

Query Result Caching

Cache query results to avoid redundant computation:

```
from iceframe.cache import QueryCache

# In-memory cache
cache = QueryCache(max_size=100)

# Use with queries
result = ice.query("users").filter(Column("age") > 30).cache(ttl=3600).execute()
```

Install: No additional dependencies required

Parallel Table Operations

Read multiple tables concurrently:

```
from iceframe.parallel import ParallelExecutor

executor = ParallelExecutor(max_workers=4)
results = executor.read_tables_parallel(ice, ["users", "orders", "products"])
```

Install: No additional dependencies required

Connection Pooling

Reuse catalog connections for better performance:

```
from iceframe.pool import CatalogPool

pool = CatalogPool(catalog_config, pool_size=5)
conn = pool.get_connection()
```

```
# Use connection
pool.return_connection(conn)
```

Install: No additional dependencies required

Memory Management

Process large tables in chunks:

```
from iceframe.memory import MemoryManager

manager = MemoryManager(max_memory_mb=1000)

# Read in chunks
for chunk in manager.read_table_chunked(ice, "huge_table", chunk_size=10000):
    process(chunk)
```

Install: `pip install "iceframe[monitoring]"` (for psutil)

Query Optimization

Automatic query optimization:

```
from iceframe.optimizer import QueryOptimizer

optimizer = QueryOptimizer()
analysis = optimizer.analyze_query("users", select_expressions, filter_expressions, group_by_expressions)
print(analysis["suggestions"])
```

Install: No additional dependencies required

Monitoring & Observability

Track query performance:

```
from iceframe.monitoring import MetricsCollector

collector = MetricsCollector()
query_id = collector.start_query("users")
# Execute query
collector.end_query(query_id, rows_returned=1000)

stats = collector.get_stats()
print(f"Avg duration: {stats['avg_duration_ms']}ms")
```

Install: `pip install "iceframe[monitoring]"` (for psutil, prometheus-client)

Streaming Support

Stream data to Iceberg tables:

```
from iceframe.streaming import StreamingWriter, stream_from_kafka

# Micro-batch streaming
writer = StreamingWriter(ice, "events", batch_size=1000)
writer.write({"id": 1, "event": "click"})
writer.flush()

# Kafka integration
stream_from_kafka(ice, "kafka-topic", "events_table", kafka_config)
```

Install: `pip install "iceframe[streaming]"` (for kafka-python)

Data Skipping

Skip unnecessary data files using statistics:

```
from iceframe.skipping import DataSkipper

skipper = DataSkipper()
stats = skipper.get_stats()
print(f"Skip rate: {stats['skip_rate']:.2%}")
```

Install: No additional dependencies required

Catalog Federation

Query across multiple catalogs:

```
from iceframe.federation import CatalogFederation

federation = CatalogFederation()
federation.add_catalog("prod", prod_config)
federation.add_catalog("dev", dev_config)

# Read from specific catalog
df = federation.read_table("prod", "users")

# Union across catalogs
combined = federation.union_tables([
    ("prod", "users"),
    ("dev", "users")
])
```

Install: No additional dependencies required

Installation

Install all scalability features:

```
pip install "iceframe[cache,streaming,monitoring]"
```

Or install individually as needed.

Source: schema_evolution.md

Schema Evolution

IceFrame provides a simple API for evolving table schemas without rewriting data.

Accessing Schema Evolution

```
# Get schema evolution interface for a table
schema_evo = ice.alter_table("users")
```

Adding Columns

```
# Add a new column
ice.alter_table("users").add_column("email", "string", doc="User email address")
```

Dropping Columns

```
# Drop a column
ice.alter_table("users").drop_column("temp_field")
```

Renaming Columns

```
# Rename a column
ice.alter_table("users").rename_column("name", "full_name")
```

Updating Column Types

Iceberg supports safe type promotion (e.g., int -> long, float -> double).

```
# Update column type
ice.alter_table("users").update_column_type("age", "long")
```

Source: statistics.md

Table Statistics

IceFrame provides comprehensive table statistics and column profiling.

Table-Level Statistics

Get overall table metadata:

```
stats = ice.stats("users")

print(f"Total snapshots: {stats['snapshots']['count']}")  
print(f"Columns: {stats['schema']['columns']}")  
print(f"Total records: {stats['data']['total_records']}")
```

Column Profiling

Profile individual columns:

```
# Numeric column
profile = ice.profile_column("users", "age")
print(f"Min: {profile['numeric_stats']['min']}")  
print(f"Max: {profile['numeric_stats']['max']}")  
print(f"Mean: {profile['numeric_stats']['mean']}")

# String column
profile = ice.profile_column("users", "name")
print(f"Avg length: {profile['string_stats']['avg_length']}")  
print(f"Distinct values: {profile['distinct_count']}")
```

Source: streaming_compaction.md

Streaming Auto-Compaction

IceFrame's `StreamingWriter` can automatically compact small files created during streaming ingestion, preventing the "small file problem" that degrades read performance.

Usage

Enable auto-compaction on the writer by calling `enable_auto_compaction`.

```
from iceframe.streaming import StreamingWriter

writer = StreamingWriter(ice, "my_table", batch_size=1000)
```

```
# Run compaction (bin-packing) after every 10 flushes
writer.enable_auto_compaction(every_n_flushes=10)

# Write data...
for record in stream:
    writer.write(record)

writer.close()
```

How it Works

When enabled, the writer tracks the number of flushes (writes to Iceberg). Once the threshold `every_n_flushes` is reached, it triggers a `bin_pack` compaction job on the table using `iceframe.compaction`. This consolidates the small data files into larger, more efficient files.

Requirements

This feature relies on the `iceframe.compaction` module. Ensure your environment supports compaction (which typically requires Spark or a compatible engine, though IceFrame's native compaction uses Pylceberg's rewrite_data_files when available).

Source: updating_tables.md

Updating Tables

IceFrame supports appending to and overwriting Iceberg tables.

Appending Data

Add new rows to an existing table.

```
import polars as pl

new_data = pl.DataFrame({
    "id": [3, 4],
    "name": ["Charlie", "David"]
})

ice.append_to_table("users", new_data)
```

You can also append using PyArrow Tables or Python dictionaries:

```
data_dict = {
    "id": [5],
    "name": ["Eve"]
}
ice.append_to_table("users", data_dict)
```

[!IMPORTANT]

Ensure data types match the table schema exactly. For example, use `int32` for `int` columns and `int64` for `long` columns.

Overwriting Data

Replace all data in the table with new data.

```
# Replaces entire table content
ice.overwrite_table("daily_report", today_data)
```

Upserts / Merge

Currently, IceFrame supports Append and Overwrite. Full Merge/Upsert functionality (Merge-on-Read) depends on the underlying Pylceberg support for your specific catalog and table version (v2).

For basic updates, you can:

Read the table

Modify the DataFrame locally

Overwrite the table (for small tables)

For large tables, use SQL-based engines (like Spark, Trino, or Dremio) connected to the same catalog for complex merge operations.

Source: variables.md

Environment Variables

IceFrame uses environment variables for configuration, authentication, and AI agent settings. You can set these in your shell or in a ` `.env` file in your project root.

Catalog Configuration

These variables configure the connection to your Iceberg catalog.

Variable	Description	Required	Default
`ICEBERG_CATALOG_URI`	The URI of your Iceberg catalog (e.g., `https://catalog.dremio.cloud/api/iceberg`).	Yes	-
`ICEBERG_CATALOG_TYPE`	The type of catalog to use.	No	`rest`
`ICEBERG_WAREHOUSE`	The warehouse location (e.g., `s3://my-bucket/warehouse`).	No	-
`ICEBERG_TOKEN`	Bearer token for authentication (alias: `ICEBERG_CATALOG_TOKEN`).	No	-
`ICEBERG_OAUTH2_SERVER_URI`	URI for the OAuth2 server if using OAuth.	No	-
`ICEBERG_CREDENTIAL`	Credential for catalog authentication (used in MCP).	No	-

`ICEBERG_CREDENTIAL_VENDING` Value for `X-Iceberg-Access-Delegation` header if using credential vending. No -

Custom Headers

You can pass custom headers to the catalog request by prefixing environment variables with `ICEBERG_HEADER_`. Underscores in the key are replaced with hyphens.

Example:

`ICEBERG_HEADER_X_Custom_Auth=secret` becomes header `X-Custom-Auth: secret`.

Dependency Environment Variables

IceFrame relies on several underlying libraries that may use their own environment variables, particularly for cloud storage and authentication.

Cloud Storage (AWS / S3)

Used by `pyiceberg`, `s3fs`, `deltalake`, `pyarrow`.

Variable	Description
:--- :---	
`AWS_ACCESS_KEY_ID` AWS Access Key.	
`AWS_SECRET_ACCESS_KEY` AWS Secret Key.	
`AWS_SESSION_TOKEN` AWS Session Token (for temporary credentials).	
`AWS_REGION` AWS Region (e.g., `us-east-1`).	
`AWS_PROFILE` Name of the AWS profile to use from `~/.aws/credentials` .	
`AWS_ENDPOINT_URL` Custom endpoint URL (useful for MinIO or S3-compatible services).	

Cloud Storage (Google Cloud / GCS)

Used by `gcsfs`, `pyiceberg`, `deltalake`.

Variable	Description
:--- :---	
`GOOGLE_APPLICATION_CREDENTIALS` Path to the JSON key file for the service account.	

Cloud Storage (Azure)

Used by `adlfs`, `pyiceberg`, `deltalake`.

Variable	Description
:--- :---	
`AZURE_STORAGE_CONNECTION_STRING` Connection string for the storage account.	
`AZURE_STORAGE_ACCOUNT_NAME` Storage account name.	
`AZURE_STORAGE_ACCOUNT_KEY` Storage account key.	
`AZURE_CLIENT_ID` Client ID for service principal.	
`AZURE_CLIENT_SECRET` Client Secret for service principal.	

| `AZURE_TENANT_ID` | Tenant ID for service principal. |

HuggingFace

Used by `datasets`.

Variable	Description
`HF_TOKEN`	Authentication token for private datasets or higher rate limits.
`HF_HOME`	Directory where datasets and models are cached (default: `~/.cache/huggingface`).

HTTP Proxies

Used by `requests` and most other network libraries.

Variable	Description
`HTTP_PROXY`	Proxy URL for HTTP requests.
`HTTPS_PROXY`	Proxy URL for HTTPS requests.
`NO_PROXY`	Comma-separated list of hosts to bypass the proxy.

AI Agent Configuration

These variables configure the AI Agent and LLM provider.

Variable	Description	Default
`ICEFRAME_LLM_PROVIDER`	Explicitly set the LLM provider (`openai`, `anthropic`, `gemini`). If not set, it is auto-detected from API keys.	Auto-detect
`ICEFRAME_LLM_MODEL`	Explicitly set the model name to use.	Provider default
`OPENAI_API_KEY`	API key for OpenAI.	-
`ANTHROPIC_API_KEY`	API key for Anthropic.	-
`GOOGLE_API_KEY`	API key for Google Gemini (alias: `GEMINI_API_KEY`).	-

Example .env File

```
# Catalog
ICEBERG_CATALOG_URI=https://catalog.example.com
ICEBERG_TOKEN=your-catalog-token
ICEBERG_WAREHOUSE=s3://my-datalake/warehouse

# AI Agent
OPENAI_API_KEY=sk-...
```

Source: visualization.md

Visualization

IceFrame integrates with **Altair** to provide declarative statistical visualization directly from your Iceberg tables.

Installation

```
pip install "iceframe[viz]"
```

Usage

Access the visualizer via the `viz` property.

```
# Plot distribution (Histogram)
chart = ice.viz.plot_distribution("my_table", "age")
chart.save("age_dist.html")

# Scatter Plot
chart = ice.viz.plot_scatter("my_table", x="age", y="salary", color="department")
chart.save("scatter.html")

# Bar Chart
chart = ice.viz.plot_bar("my_table", x="department", y="salary")
chart.save("bar.html")

# Line Chart
chart = ice.viz.plot_line("my_table", x="date", y="sales")
chart.save("sales.html")
```

Performance Note

The visualizer automatically limits the number of rows fetched (default 10,000) to prevent crashing your browser or notebook. You can adjust this limit via the `limit` parameter, but be cautious with large datasets. For larger data, consider aggregating using `ice.query()` first and then plotting the result.